

# SuiteSparse:GraphBLAS: a parallel implementation of the GraphBLAS specification

Tim Davis, Texas A&M University

June 29, 2020

SIAM MDS'20: Linear Algebraic Tools for Graph Computation

# SuiteSparse:GraphBLAS matrix data structure

- both CSC (compressed sparse column form) and CSR. Default is CSR.
- sparse ( $O(n + |A|)$  space): a dense vector of sparse vectors
- or hypersparse ( $O(|A|)$ ): a sparse vector of sparse vectors
- **ZOMBIES!**: an entry marked for deletion (negated row index so binary search still OK)
- *pending tuples*: an entry in an ordered list, waiting to be inserted
- values: just about anything (each scalar of fixed size)
- no value given to the implicit entry

## Parallel algorithms (assume CSC)

- matrix-matrix and matrix-vector multiply. Three primary methods:
  - saxpy-style (Gustavson and hash)
  - hash method (Nagasaka, Matsuoka, Azad, Buluç (2018); extended here with a concurrent data structure via atomics)
  - dot, with mask:  $\mathbf{C}\langle\mathbf{M}\rangle = \mathbf{AB}'$
  - dot, without mask but  $\mathbf{C}$  dense (in-place)
- dot-style with mask: each  $M_{ij}$
- saxpy-style: 4 kinds of tasks, all can be used in a single  $\mathbf{C} = \mathbf{AB}$ . Tasks selected based on amount of work per column, and number of threads. Let  $f$  be the flop count.
- Parallel assignment:  $\mathbf{C}(\mathbf{I}, \mathbf{J}) = \mathbf{A}$ , with **ZOMBIES!** ... and pending tuples too.

## Parallel matrix multiply (assume CSC)

Four kinds of saxpy tasks, each with 3 variants: no mask,  $M$ , and  $\neg M$ .

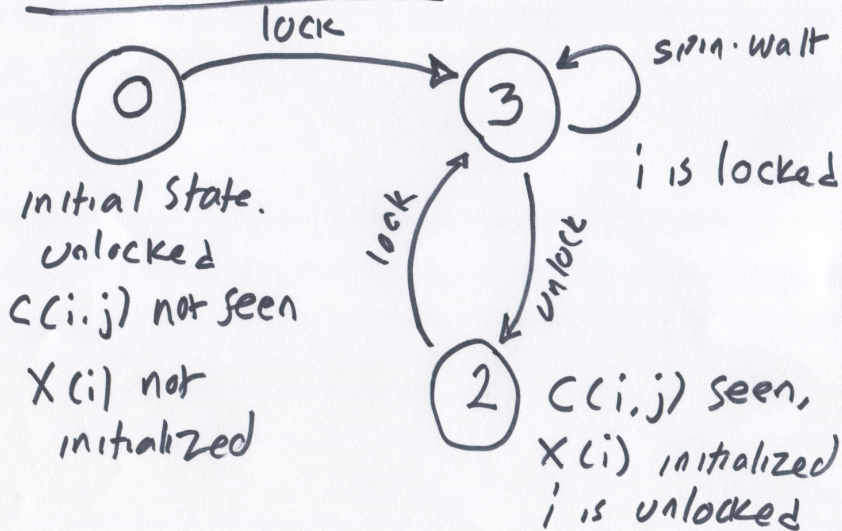
- coarse Gustavson task:  $C(:, j1:j2) = A * B(:, j1:j2)$  by one thread with  $O(n)$  workspace, owns columns  $j1:j2$ . Time is  $O(n + f)$ .
- fine Gustavson task: multiple threads cooperate on a single column:  $C(:, j) = A * B(:, j)$ . Threads share a single  $O(n)$  workspace; using atomics. Time is  $O(n/p + f/p)$ .
- coarse hash task:  $C(:, j1:j2) = A * B(:, j1:j2)$ , by one thread using  $O(h)$  workspace,  $h \ll n$ , where  $h = 4f$ . Time is  $O(f)$  assuming few collisions.
- fine hash task: multiple threads cooperate on a single column:  $C(:, j) = A * B(:, j)$ . Threads share  $O(h)$  workspace,  $h \ll n$ , where  $h = 4f$ . Time is  $O(f/p)$  assuming few collisions.

## Fine Gustavson task: with no mask

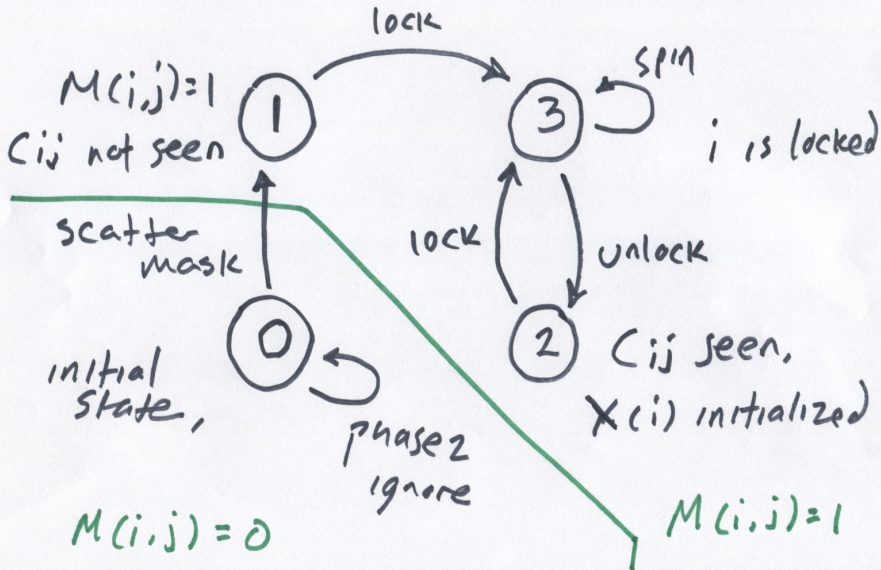
To compute both pattern and values of  $C$ .

- concurrent data structure: two arrays of size  $n$ :  $X$  for numerical values,  $F$  for state (`int8`) for  $n$  concurrent finite-state machines.  $F$  is calloc'd so starts as zero.
- phase1: scatter the mask (no mask for this case however).
- phase2: scatter/sum values in  $X$ , using machine  $F$ . Using atomics. Time is  $O(f)$  assuming no spin-wait.
- phase3: count  $c(j) = \text{nnz}(C(:,j))$  for each  $j$ . Time is  $O(n/p)$  with  $p$  threads for a single column  $j$ .
- phase4:  $C_p = \text{cumsum}(c)$  for column pointers, for all  $j$  Time is  $O(n/p)$  for entire matrix.
- phase5: gather from  $X$  to create values and pattern of  $C(:,j)$ .

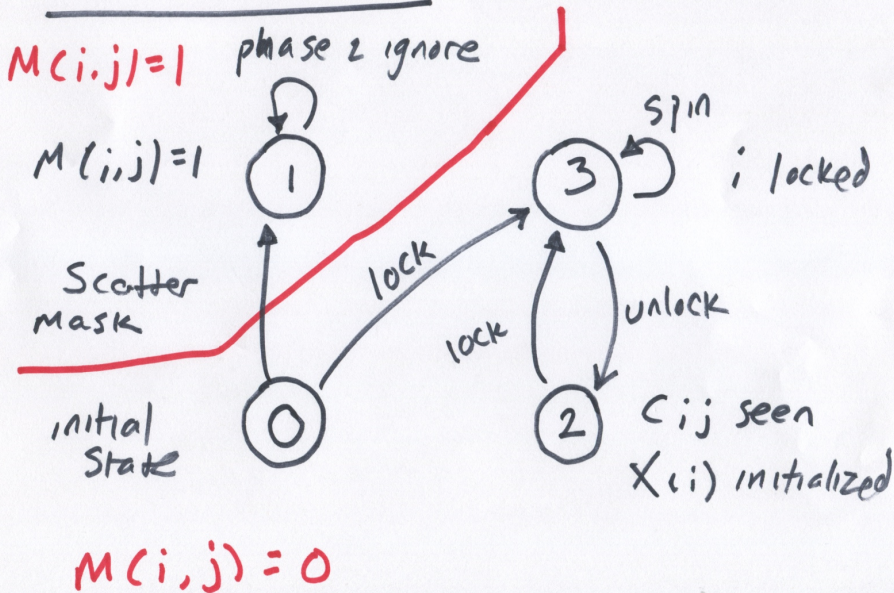
# Fine Gustavson task : no mask



# Fine Gustafson task: with mask M



# Fine Gustavson task : with !M





## Fine hash tasks

- $F$  has size  $4f$ , 64-bit integers:  $F[k] = (h, f)$ .
- $f$ : lowest 2 bits: for 4-state finite-state machine
- $h$ : upper 62 bits:  $h = \text{row index } i + 1$  occupying the hash entry  $F[k]$
- $F[k] = (h, f)$  can be read/swapped/written in a single atomic operation
- The four states:
  - $h=0, f=0$ : unlocked, unoccupied.
  - $h=i+1, f=1$ : unlocked, occupied by row  $i$ .  $C(i, j)$  not seen, or ignored.  $X(k)$  not initialized.
  - $h=i+1, f=2$ : unlocked, occupied by row  $i$ .  $C(i, j)$  seen.  $X(k)$  initialized.
  - $h=\text{anything}, f=3$ : locked, occupied by something.
- simple hash function; linear probing if first entry occupied.
- hash table size based on flop count  $f$  for that column  $j$ .

## Parallel matrix-matrix performance

$C=A*A$  for matrix ND/nd3k ( $n = 9000$ , with 3.3 million entries): All results on 20-core Intel Xeon. MATLAB R2018a.

time in seconds:

method	single	double	single complex	double complex
MATLAB	-	3.76	-	7.90
GrB:1	3.12	3.35	4.63	4.92
GrB:20	0.22	0.24	0.33	0.31
GrB:40	0.18	0.21	0.27	0.29

speedup:

vs GrB:1	17.3	16.0	17.1	17.0
vs MATLAB	20.9	18.0	29.3	27.3

## Parallel matrix - sparse vector performance

$C=A*x$  for Freescale2 ( $n = 3$  million, 14.3 million entries).  $x$  is 50% nonzero:

time in seconds:

method	single	double	single complex	double complex
MATLAB	-	0.228	-	0.288
GrB:1	0.090	0.096	0.121	0.174
GrB:20	0.010	0.014	0.015	0.024
GrB:40	0.011	0.014	0.015	0.029

speedup:

vs GrB:1	8.2	6.9	8.1	6.0
vs MATLAB	20.7	16.3	19.2	9.9

# Parallel $C(I,J)=A$

- 128 variants:  $C(i,j)\langle M \rangle = C(i,j) \odot A$ 
  - **M** present, or not
  - mask complemented, or not
  - mask structural, or not
  - REPLACE option enabled, or not
  - accumulator  $\odot$  present, or not
  - **A**: scalar or matrix
  - **S** matrix constructed: see below
- implemented with 30 main kernels, including 7 special cases:
  - $C=x$  with scalar  $x$
  - $C=A$
  - $C+=x$  when  $C$  is dense
  - $C+=A$  when  $C$  is dense
  - $C\langle M \rangle=x$ , when  $C$  is dense
  - $C\langle A \rangle=A$ , when  $C$  is dense
  - $C\langle M, \text{struct} \rangle=A$ , when  $A$  is dense and  $C$  is empty

## Parallel $C(I,J)=A$ , simple case

Basic case: no mask, no accumulator, no replace option, mask not complemented, I and J arbitrary lists of indices, C and A sparse, ...

- 1st pass: structural extraction to compute S
  - extract  $S=C(I,J)$  where  $S(i,j)$  is the not *value* of  $C(I[i],J[j])$ , but its *position* in the data structure for C.
  - S and A have the same size.
  - $S(i,j)$  tells where  $A(i,j)$  goes. Let  $p=S(i,j)$
- 2nd pass: do updates, make zombies, count pending tuples
  - $S(i,j)$  and  $A(i,j)$  both present. update:  $C(p) = A(i,j)$ . A **ZOMBIE** might come back to life!
  - $S(i,j)$  not present;  $A(i,j)$  present.  $A(i,j)$  becomes a *pending tuple*, waiting to be inserted. Lazy.
  - $S(i,j)$  present;  $A(i,j)$  *not* present. Entry at  $C(p)$  must be deleted ... becomes a **ZOMBIE**!
- cumulative sum of pending tuples found per thread
- 3rd pass: each thread puts its pending tuples in common list

## Parallel $C=A(I,J)$ performance

A is square,  $n=100$  million, 1 billion entries.

$C = A(I,J)$ , where  $I=\text{randperm}(n,n/10)$ , also J.

Note: GrB algorithm not presented in this talk. Results here to compare with next slide on  $C(I,J)=A$ .

	threads	time (sec)	speedup vs MATLAB	speedup vs GrB:1
MATLAB	1	8.87	1	1.07
GrB	1	9.48	0.94	1
GrB	5	2.70	3.28	3.51
GrB	10	1.90	4.66	4.98
GrB	20	1.48	5.98	6.39
GrB	40	1.35	6.56	7.01

## Parallel $C(I,J)=A$ performance

Same A, I, and J as last slide.

$A(I,J) = 2*A(I,J)$  (no change to pattern).

	threads	time (sec)	speedup vs MATLAB	speedup vs GrB:1
MATLAB	1	> 24 hours	-	-
GrB	1	32.02	> 2,700	1
GrB	5	8.88	> 10,000	3.60
GrB	10	5.91	> 15,000	5.42
GrB	20	4.78	> 18,000	6.69
GrB	40	4.43	> 20,000	7.32

GrB: about 3x the time for  $C=A(I,J)$  but this expression starts with that; remainder is  $A(I,J)=2*C$ .

MATLAB: still running after 24+ hours. GrB using same syntax, via MATLAB @GrB interface.

## Parallel $C(I,J)=A$ performance

Same A, I, and J as last slide.

$A(I, J) = 2 * A(I+1, J+1)$  (changes pattern).

	threads	time (sec)	speedup vs MATLAB	speedup vs GrB:1
MATLAB	1	-	-	0
GrB	1	55.25	-	1
GrB	5	14.79	-	3.74
GrB	10	9.59	-	5.76
GrB	20	7.47	-	7.39
GrB	40	7.04	-	7.84

MATLAB: first experiment still running after 24 hours.



## Parallel performance: betweenness centrality

Dominated by matrix-matrix multiply (one matrix 4-by- $n$  and dense)

time in seconds

matrix	threads			
	1	10	20	40
kron	1076.9	137.7	74.6	42.3
urand	1405.8	107.2	70.3	63.2
twitter	328.1	35.1	17.5	13.0
web	91.3	13.4	8.3	7.8
road	52.5	51.9	54.1	61.4

speedup

kron	1	7.8	14.4	25.5
urand	1	13.1	20.0	22.2
twitter	1	9.3	18.7	25.2
web	1	6.8	11.0	11.7
road	1	1.0	1.0	0.9

## Parallel performance: pagerank

Dominated by matrix-vector multiply (dense vector).  
time in seconds

matrix	threads			
	1	10	20	40
kron	372.7	40.8	22.2	21.8
urand	378.9	42.2	27.7	27.8
twitter	286.5	32.0	17.9	17.3
web	90.0	10.9	8.9	8.9
road	12.9	1.8	1.4	1.4

speedup

kron	1	9.1	16.8	17.1
urand	1	9.0	13.7	13.6
twitter	1	9.0	16.0	16.6
web	1	8.3	10.1	10.1
road	1	7.2	9.2	9.2