

SuiteSparse:GraphBLAS: a parallel implementation of the GraphBLAS specification

Tim Davis, Texas A&M University

June 29, 2020

SIAM MDS'20: Linear Algebraic Tools for Graph Computation

SuiteSparse:GraphBLAS matrix data structure

- both CSC (compressed sparse column form) and CSR. Default is CSR.
- sparse ($O(n + |A|)$ space): a dense vector of sparse vectors
- or hypersparse ($O(|A|)$): a sparse vector of sparse vectors
- **ZOMBIES!**: an entry marked for deletion (negated row index so binary search still OK)
- *pending tuples*: an entry in an ordered list, waiting to be inserted
- values: just about anything (each scalar of fixed size)
- no value given to the implicit entry

Parallel algorithms (assume CSC)

- matrix-matrix and matrix-vector multiply. Three primary methods:
 - saxpy-style (Gustavson and hash)
 - hash method (Nagasaka, Matsuoka, Azad, Buluç (2018); extended here with a concurrent data structure via atomics)
 - dot, with mask: $\mathbf{C}\langle\mathbf{M}\rangle = \mathbf{AB}'$
 - dot, without mask but \mathbf{C} dense (in-place)
- dot-style with mask: each M_{ij}
- saxpy-style: 4 kinds of tasks, all can be used in a single $\mathbf{C} = \mathbf{AB}$. Tasks selected based on amount of work per column, and number of threads. Let f be the flop count.
- Parallel assignment: $\mathbf{C}(\mathbf{I}, \mathbf{J}) = \mathbf{A}$, with **ZOMBIES!** ... and pending tuples too.

Parallel matrix multiply (assume CSC)

Four kinds of saxpy tasks, each with 3 variants: no mask, M , and $\neg M$.

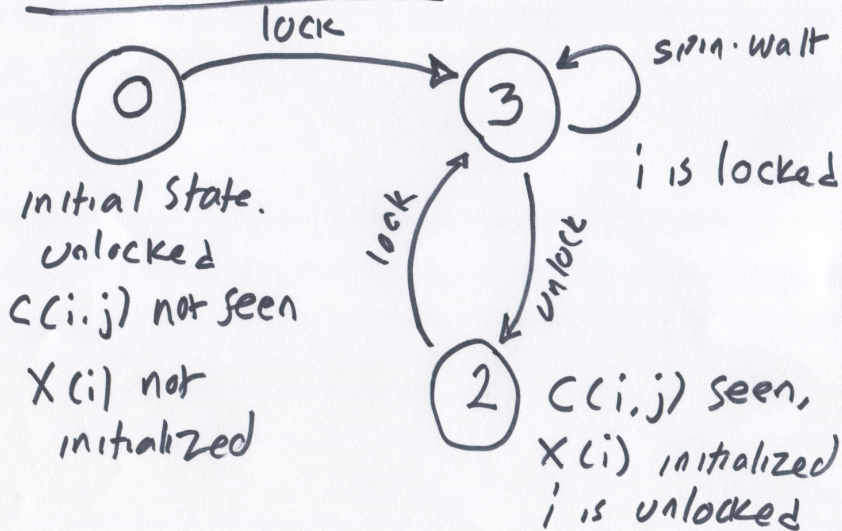
- coarse Gustavson task: $C(:, j1:j2) = A * B(:, j1:j2)$ by one thread with $O(n)$ workspace, owns columns $j1:j2$. Time is $O(n + f)$.
- fine Gustavson task: multiple threads cooperate on a single column: $C(:, j) = A * B(:, j)$. Threads share a single $O(n)$ workspace; using atomics. Time is $O(n/p + f/p)$.
- coarse hash task: $C(:, j1:j2) = A * B(:, j1:j2)$, by one thread using $O(h)$ workspace, $h \ll n$, where $h = 4f$. Time is $O(f)$ assuming few collisions.
- fine hash task: multiple threads cooperate on a single column: $C(:, j) = A * B(:, j)$. Threads share $O(h)$ workspace, $h \ll n$, where $h = 4f$. Time is $O(f/p)$ assuming few collisions.

Fine Gustavson task: with no mask

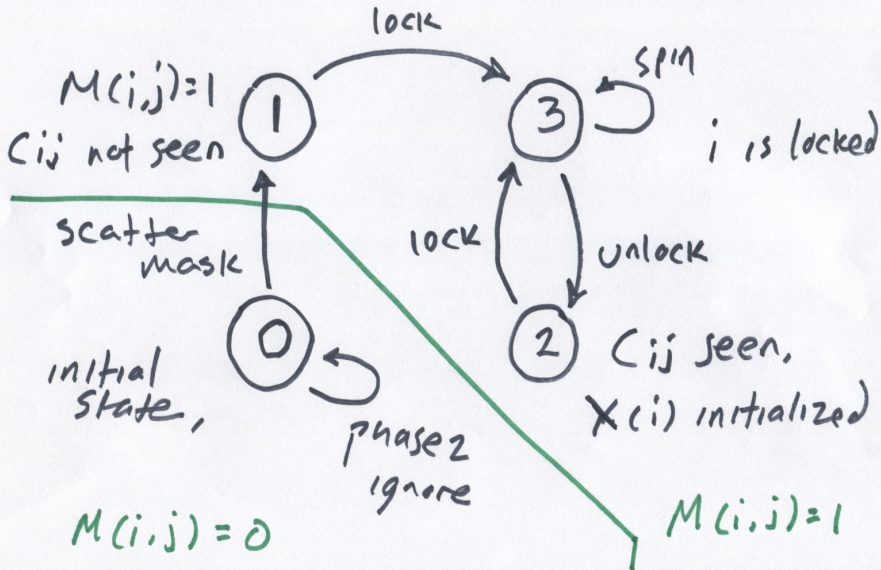
To compute both pattern and values of C .

- concurrent data structure: two arrays of size n : X for numerical values, F for state (`int8`) for n concurrent finite-state machines. F is calloc'd so starts as zero.
- phase1: scatter the mask (no mask for this case however).
- phase2: scatter/sum values in X , using machine F . Using atomics. Time is $O(f)$ assuming no spin-wait.
- phase3: count $c(j) = \text{nnz}(C(:,j))$ for each j . Time is $O(n/p)$ with p threads for a single column j .
- phase4: $C_p = \text{cumsum}(c)$ for column pointers, for all j Time is $O(n/p)$ for entire matrix.
- phase5: gather from X to create values and pattern of $C(:,j)$.

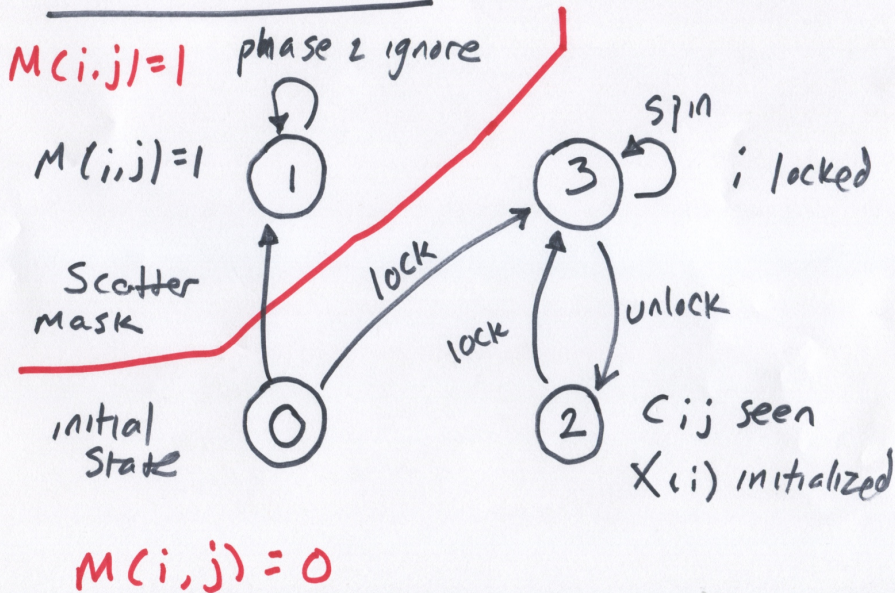
Fine Gustavson task : no mask



Fine Gustafson task: with mask M



Fine Gustavson task : with !M



Fine hash tasks

- F has size $4f$, 64-bit integers: $F[k] = (h, f)$.
- f : lowest 2 bits: for 4-state finite-state machine
- h : upper 62 bits: $h = \text{row index } i + 1$ occupying the hash entry $F[k]$
- $F[k] = (h, f)$ can be read/swapped/written in a single atomic operation
- The four states:
 - $h=0, f=0$: unlocked, unoccupied.
 - $h=i+1, f=1$: unlocked, occupied by row i . $C(i, j)$ not seen, or ignored. $X(k)$ not initialized.
 - $h=i+1, f=2$: unlocked, occupied by row i . $C(i, j)$ seen. $X(k)$ initialized.
 - $h=\text{anything}, f=3$: locked, occupied by something.
- simple hash function; linear probing if first entry occupied.
- hash table size based on flop count f for that column j .

Parallel matrix-matrix performance

$C=A*A$ for matrix ND/nd3k ($n = 9000$, with 3.3 million entries): All results on 20-core Intel Xeon. MATLAB R2018a.

time in seconds:

method	single	double	single complex	double complex
MATLAB	-	3.76	-	7.90
GrB:1	3.12	3.35	4.63	4.92
GrB:20	0.22	0.24	0.33	0.31
GrB:40	0.18	0.21	0.27	0.29

speedup:

vs GrB:1	17.3	16.0	17.1	17.0
vs MATLAB	20.9	18.0	29.3	27.3

Parallel matrix - sparse vector performance

$C=A*x$ for Freescale2 ($n = 3$ million, 14.3 million entries). x is 50% nonzero:

time in seconds:

method	single	double	single complex	double complex
MATLAB	-	0.228	-	0.288
GrB:1	0.090	0.096	0.121	0.174
GrB:20	0.010	0.014	0.015	0.024
GrB:40	0.011	0.014	0.015	0.029

speedup:

vs GrB:1	8.2	6.9	8.1	6.0
vs MATLAB	20.7	16.3	19.2	9.9

Parallel $C(I,J)=A$

- 128 variants: $C(i,j)\langle M \rangle = C(i,j) \odot A$
 - **M** present, or not
 - mask complemented, or not
 - mask structural, or not
 - REPLACE option enabled, or not
 - accumulator \odot present, or not
 - **A**: scalar or matrix
 - **S** matrix constructed: see below
- implemented with 30 main kernels, including 7 special cases:
 - $C=x$ with scalar x
 - $C=A$
 - $C+=x$ when C is dense
 - $C+=A$ when C is dense
 - $C\langle M \rangle=x$, when C is dense
 - $C\langle A \rangle=A$, when C is dense
 - $C\langle M, \text{struct} \rangle=A$, when A is dense and C is empty

Parallel $C(I,J)=A$, simple case

Basic case: no mask, no accumulator, no replace option, mask not complemented, I and J arbitrary lists of indices, C and A sparse, ...

- 1st pass: structural extraction to compute S
 - extract $S=C(I,J)$ where $S(i,j)$ is the not *value* of $C(I[i],J[j])$, but its *position* in the data structure for C.
 - S and A have the same size.
 - $S(i,j)$ tells where $A(i,j)$ goes. Let $p=S(i,j)$
- 2nd pass: do updates, make zombies, count pending tuples
 - $S(i,j)$ and $A(i,j)$ both present. update: $C(p) = A(i,j)$. A **ZOMBIE** might come back to life!
 - $S(i,j)$ not present; $A(i,j)$ present. $A(i,j)$ becomes a *pending tuple*, waiting to be inserted. Lazy.
 - $S(i,j)$ present; $A(i,j)$ *not* present. Entry at $C(p)$ must be deleted ... becomes a **ZOMBIE**!
- cumulative sum of pending tuples found per thread
- 3rd pass: each thread puts its pending tuples in common list

Parallel $C=A(I,J)$ performance

A is square, $n=100$ million, 1 billion entries.

$C = A(I,J)$, where $I=\text{randperm}(n,n/10)$, also J.

Note: GrB algorithm not presented in this talk. Results here to compare with next slide on $C(I,J)=A$.

	threads	time (sec)	speedup vs MATLAB	speedup vs GrB:1
MATLAB	1	8.87	1	1.07
GrB	1	9.48	0.94	1
GrB	5	2.70	3.28	3.51
GrB	10	1.90	4.66	4.98
GrB	20	1.48	5.98	6.39
GrB	40	1.35	6.56	7.01

Parallel $C(I,J)=A$ performance

Same A, I, and J as last slide.

$A(I,J) = 2*A(I,J)$ (no change to pattern).

	threads	time (sec)	speedup vs MATLAB	speedup vs GrB:1
MATLAB	1	> 24 hours	-	-
GrB	1	32.02	> 2,700	1
GrB	5	8.88	> 10,000	3.60
GrB	10	5.91	> 15,000	5.42
GrB	20	4.78	> 18,000	6.69
GrB	40	4.43	> 20,000	7.32

GrB: about 3x the time for $C=A(I,J)$ but this expression starts with that; remainder is $A(I,J)=2*C$.

MATLAB: still running after 24+ hours. GrB using same syntax, via MATLAB @GrB interface.

Parallel $C(I,J)=A$ performance

Same A, I, and J as last slide.

$A(I, J) = 2 * A(I+1, J+1)$ (changes pattern).

	threads	time (sec)	speedup vs MATLAB	speedup vs GrB:1
MATLAB	1	-	-	0
GrB	1	55.25	-	1
GrB	5	14.79	-	3.74
GrB	10	9.59	-	5.76
GrB	20	7.47	-	7.39
GrB	40	7.04	-	7.84

MATLAB: first experiment still running after 24 hours.

Parallel performance: betweenness centrality

Dominated by matrix-matrix multiply (one matrix 4-by- n and dense)

time in seconds

matrix	threads			
	1	10	20	40
kron	1076.9	137.7	74.6	42.3
urand	1405.8	107.2	70.3	63.2
twitter	328.1	35.1	17.5	13.0
web	91.3	13.4	8.3	7.8
road	52.5	51.9	54.1	61.4

speedup

kron	1	7.8	14.4	25.5
urand	1	13.1	20.0	22.2
twitter	1	9.3	18.7	25.2
web	1	6.8	11.0	11.7
road	1	1.0	1.0	0.9

Parallel performance: pagerank

Dominated by matrix-vector multiply (dense vector).
time in seconds

matrix	threads			
	1	10	20	40
kron	372.7	40.8	22.2	21.8
urand	378.9	42.2	27.7	27.8
twitter	286.5	32.0	17.9	17.3
web	90.0	10.9	8.9	8.9
road	12.9	1.8	1.4	1.4

speedup

kron	1	9.1	16.8	17.1
urand	1	9.0	13.7	13.6
twitter	1	9.0	16.0	16.6
web	1	8.3	10.1	10.1
road	1	7.2	9.2	9.2

MATLAB vs GraphBLAS

GraphBLAS is often many times faster than MATLAB, even when using the same syntax via operator overloading. **However:**

- GraphBLAS would not exist without MATLAB. GraphBLAS is complex, and to test it, I wrote it twice: in C and in simple MATLAB (3 nested loops for $C=A*B$ for example), with dense matrices and intentionally slow. The MATLAB code is **not** tested for performance (not here, not ever). The advantage is that it so clear that reads like the spec. GraphBLAS has a *very* complex spec (230+ pages).
- MATLAB is not intrinsically slower. It can be much faster; its sparse matrix operations for $C=A*B$, $C(I,J)=A$, etc, could be done via calls to GraphBLAS.
- Some operations are still faster in MATLAB: $C = [A \ B]$ for example, is 3x faster in MATLAB vs GraphBLAS, since I have yet to write a fast method in GraphBLAS for concatenation.